

Compare and Contrast Maximum Likelihood Method and Inverse Probability Weighting Method in Missing Data Analysis

Scott Sun

Advisor: Dr. Eric Reyes

Abstract

Data can be lost for different reasons, but sometimes the missingness is a part of the data collection process. Unbiased and efficient estimation of the parameters governing the response mean model requires the missing data to be appropriately addressed. This paper compares and contrasts the Maximum Likelihood and Inverse Probability Weighting estimators in an Outcome-Dependent Sampling design that deliberately generates incomplete observations. We demonstrate the comparison through numerical simulations under varied conditions: different coefficients of determination, and whether or not the mean model is misspecified.

Keyword: Outcome-Dependent Sampling, Maximum Likelihood, Inverse Probability Weighting, EM Algorithm



Rose-Hulman Institute of Technology
Spring, 2021

1 Introduction

Consider a hypothetical study about the association between oxygen levels and the age of patients with COVID-19. Researchers may conduct it in two stages, sampling patients' ages and oxygen levels in the first stage. Based on the information obtained from the first stage, the researchers identify some elderly patients to be critically ill. Then, in the second stage, the researcher will collect detailed information such as medical history from these seriously ill patients. This hypothetical study is an example of an Outcome-Dependent Sampling (ODS) study, where data are collected in two stages. In the first stage, information is available on the response and some of the covariates for all observations. However, in the second stage, information is available on other covariates among only a subset of the sample (Zhao & Lipsitz, 1992). Further, the likelihood of observing an individual in the second stage can depend on data collected during the first stage. This sampling method allows researchers to concentrate resources on places with the most valuable information according to their topic (Weaver & Zhou, 2005). However, this unique design deliberately generates incomplete observations in the final dataset. Therefore, estimation of the parameters governing the model for the response requires the missing data to be appropriately addressed.

Among applied researchers, it is common only to retain complete cases when data is subject to missingness. This approach is called complete-case (CC) analysis, also known as case deletion and listwise deletion, and is usually the default in many software packages. It is admittedly straightforward to understand but not always justified since it may exclude potentially helpful information. In an ODS study, missingness has a systematic pattern and is intrinsic to the data collection, where the likelihood of observing a complete case varies among observations. Therefore, estimates directly obtained from the complete-case analysis can be biased and inefficient (Schafer & Graham, 2002).

For general missing data problems, three primary methods have been developed as alternatives to the complete-case approach: Multiple Imputation, Maximum Likelihood (ML), and Inverse Probability Weighting (IPW) (Schafer & Graham, 2002). Multiple Imputation solves the missingness by creating different plausible versions of the data by imputing the missing values for multiple times and aggregating the results. ML estimation obtains the estimates directly by optimizing a likelihood function that incorporates the impact of the missingness. In the IPW approach, the analysis model is fitted only to the complete observations, but different weights are assigned to adjust individual contribution to the analysis based on the missingness structure.

In this paper, we only discuss ML and IPW; the Multiple Imputation method is outside the scope of this paper. Earlier works have compared various ML estimators and IPW estimators, with the ML methods tending to outperform the IPW methods in efficiency. However, Seaman et al. (2011) discuss ways to improve IPW estimators' efficiency by truncating the weights before assigning them to the corresponding observation. The previous works did not compare the ML estimators with estimators calculated by this improved version of the IPW method. Moreover, while the earlier works theorized on the result, they did not thoroughly compare the two methods when the mean response model is misspecified. Therefore, in this article, we would like to compare the ML estimators with the weight-stabilized IPW estimators in an Outcome-Dependent Sampling design under varied conditions. We will

demonstrate the comparisons through numerical simulations.

2 Types of Missingness Mechanisms

Missingness, or incompleteness, is usually represented by a matrix of Bernoulli random variables, R . Within the scope of this paper, R is reduced as follows:

$$R_i = \begin{cases} 1, & \text{if the } i\text{-th subject is fully observed} \\ 0, & \text{otherwise} \end{cases}$$

It is dictated by the missingness mechanism, which is a part of the data generation process. The mechanism plays a crucial role in the study of missing-data analysis (Schafer & Graham, 2002). It helps to characterize the relationship between the missingness matrix and the original data matrix. There exist three types of missingness mechanisms: Missing Completely at Random (MCAR), Missing at Random (MAR), and Missing Not at Random (MNAR). In this paper, we will only discuss MCAR and MAR.

Let the data matrix $X = (x_{ij})$, which consists of X_{mis} , the missing variables, and X_{obs} , the observed variables. In a missing variable, some subjects are unobserved; however, in an observed variable, the value is observed for all observations. Moreover, let the missingness vector $R = (R_1, R_2 \dots R_i \dots)$. Mathematically, the mechanism can be classified by the conditional distribution $f(R|X, \Theta)$, where Θ is an unknown parameter vector.

2.1 Missing Completely at Random

If the likelihood of being observed is independent of any variables, the missingness mechanism is MCAR. That is,

$$\Pr(R|X, \Theta) = \Pr(R|\Theta).$$

MCAR is a strong ideal assumption on the mechanism, which assumes the missingness is unrelated to all variables. It indicates that the pattern is not affected by the studied subjects. For instance, when we research blood samples, several samples can be contaminated during the delivery process. Thus, we cannot collect information from these polluted specimens. Nothing about the specimens made them more or less likely to be contaminated.

2.2 Missing at Random

If the likelihood of being observed depends only on those fully observed variables, the missingness mechanism is MAR. That is,

$$\Pr(R|X, \Theta) = \Pr(R|X_{obs}, \Theta).$$

Compared to MCAR, MAR is a less constrained statement. Based on this definition, MCAR can be seen as an extreme, special case of MAR (Schafer & Graham, 2002). Due to the unique design of ODS, the missingness follows the pattern of MAR. The likelihood an observation is completely observed depends on data collected during the first stage, where all variables are fully observed.

3 ODS Design and Likelihood Function

In Section 1, we described the data collection of an ODS study in general. Now, we present a specific sampling process with notation and detail. In the first stage, we sample N individuals from the population and obtain \mathbf{Y} , a continuous outcome variable, and \mathbf{X} , a continuous predictor variable, from all of them. In the second stage, we obtain \mathbf{Z} , a binary categorical predictor variable, from a subset of the sample subjects. For each individual, the likelihood \mathbf{Z} is observed depends on the values for variables obtained in the first stage: \mathbf{Y} , and \mathbf{X} .

We define \mathbf{R} to indicate the completeness of the observations in the data set. Under this setting, \mathbf{R} indicates if \mathbf{Z} is obtained from the individuals. Thus, the missingness mechanism is MAR as \mathbf{R} depends only one the observed variables. Let S represent the index set of all individuals of complete observation, and let \bar{S} represent the index set of all individuals whose Z_i is missing where $i = 1, \dots, N$. That is, $S = \{i : R_i = 1\}$ and $\bar{S} = \{i : R_i = 0\}$.

Zhao et al. (1992) discuss the likelihood function for two-stage case-control designs. In their setting, the response variable is binary, and the first predictor variable is discrete. However, the ODS design for continuous outcomes is comparable to the case-control design in terms of the missingness mechanism. The general expression of the observed likelihood function is invariant to variables' types and distributions. Therefore, the observed likelihood for a data set with this ODS design is

$$L_{Obs}(\Theta) = \prod_{i=1}^N \left[f(y_i, x_i; \Theta) f(z_i | y_i, x_i; \Theta) \right]^{R_i} \left[f(y_i, x_i; \Theta) \right]^{1-R_i}$$

The goal of interest is to modify and utilize the likelihood function to perform estimation and inference for Θ , a vector of parameters that describe the variables.

4 Analysis Methods

4.1 Complete-Case Analysis

As we have introduced in Section 1, CC analysis is a common method in the presence of data missingness. After we remove the incomplete observations from the data set, all the remaining observations in the edited data set are complete. These remaining observations are indexed by elements in S . For all $i \in S$, $R_i = 1$. As a result, the likelihood function is modified to $L_{CC}(\Theta)$.

$$\begin{aligned} L_{CC}(\Theta) &= \prod_{i \in S} \left[f(y_i, x_i; \Theta) f(z_i | y_i, x_i; \Theta) \right]^{R_i} \left[f(y_i, x_i; \Theta) \right]^{1-R_i} \\ &= \prod_{i \in S} \left[f(y_i, x_i; \Theta) f(z_i | y_i, x_i; \Theta) \right] \end{aligned}$$

Then, the CC analysis obtains $\hat{\Theta}$ by maximizing $L_{CC}(\Theta)$. There are a few circumstances where the CC analysis can yield valid estimators. When the missingness mechanism is

MCAR, the estimator is generally unbiased and consistent because the completeness is independent of all variables in the data set. The complete cases can be viewed as random samples from an imagined full data set (Schafer & Graham, 2002; Seaman & White, 2011). In our specific ODS design notations, $\Pr(R = 1|X, Y; \Theta) = \Pr(R = 1)$. Second, the CC analysis estimator can have negligible bias under a weaker condition: the missingness mechanism is MAR, but it is independent of the response variable given the observed predictor variables. This assumes the mean model is correctly specified (White & Carlin, 2010; Seaman & White, 2011). In our notations, $\Pr(R = 1|X, Y; \Theta) = \Pr(R = 1|X; \Theta)$.

In situations other than the two mentioned above, the CC analysis generally gives biased estimators. The ODS study design does not belong to the two scenarios since the missingness mechanism of MAR and the case completeness variable, R_i , depends on the outcome variable, Y_i , for all observations by definition. Therefore, the CC analysis will not be a preferable approach to address the data missingness for data sets generated using the ODS scheme.

4.2 Maximum Likelihood

The ML approach does not edit the original data set. The analysis model is fitted to all the observations in the data set. The essential step in the ML method is to construct an appropriate likelihood function that captures all the available information from it. In this paper, the ML estimator we discuss is fully parametric, obtained by maximizing the observed likelihood function $L_{Obs}(\Theta)$. To perform maximum likelihood estimation, we have to posit models on the distributions of the response variable, $f(Y|X, Z)$, and the predictor variables, $f(X, Z)$.

The EM algorithm, an iterative computation algorithm, is a common approach to calculate the maximum likelihood estimator since it is difficult to derive the expression of the estimator in a closed-form. A detailed computing procedure for our ODS design is included in the Appendix. The observed information matrix of the estimate is estimated by

$$I_n(\hat{\Theta}) = n \cdot \text{Cov}(\mathbf{y}, \mathbf{x}, \mathbf{z}, \hat{\Theta})$$

The standard error of $\hat{\Theta}_i$ is calculated by the square root of the i -th diagonal element of $I_n(\hat{\Theta})^{-1}$.

However, the limitation of this parametric method is that it heavily relies on assumptions. If correct assumptions are provided, the ML estimators are efficient and fully utilizing the information from both stages (Zhao & Lipsitz, 1992). However, if some assumptions are incorrect (for instance, the response model is misspecified), this method has poor performance estimation and inference. Moreover, the EM algorithm is case-sensitive since it is closely related to the specific data set and its likelihood function. Thus, we have to develop a different corresponding EM algorithm for a data set with a dissimilar design and likelihood function.

Weaver et al. (2005) discuss semiparametric ML estimators, such as the Maximum Semiparametric Empirical Likelihood Estimator and the Maximum Estimated Likelihood Estimator, for ODS studies. The predictors' probability density functions are replaced by nonparametric estimated density functions in the observed likelihood functions in these methods. With

fewer assumptions, the new likelihood functions can solve problems from a broader class of designs, where the predictor variables can be from any distribution.

4.3 Inverse Probability Weighting

The IPW approach, unlike the parametric ML method, requires fewer assumptions: it only posits models for the mean of the response variable and the missingness mechanism. Different from the complete case analysis, the IPW method weights each individual in the complete sample by the inverse probability of being fully observed (Seaman & White, 2011). Instead of calculating $\hat{\Theta}$ by optimizing the log-likelihood $l_{Obs}(\Theta)$, we obtain it by optimizing a weighted log-likelihood function (Weaver & Zhou, 2005). In our ODS design, the weighted log-likelihood function is

$$l_{IPW}(\Theta) = \sum_{i=1}^N \frac{1}{p_i} R_i \cdot \log f(Y_i | X_i, Z_i; \Theta)$$

where p_i is the probability to have the i -th observation be fully observed.

The probability p_i is estimated by using logistic regression. The logistic regression models the probability of complete observation given all the observed variables, $\Pr(R_i | X_i, Y_i)$ in our case. Weaver et al. (2005) prove that using the observed or estimated probability for weighting can generate more efficient estimators than using the known probability.

IPW, as a semiparametric method, is less constrained by assumptions for the predictor variables so that we can apply it to a broader range of problems. However, the estimator can be less efficient than the ML estimator when the model is well-specified (Schafer & Graham, 2002). The inefficiency is reflected by the IPW estimates' large standard errors. These large standard errors are generated as the complete cases can be assigned with large weights. Some people argue that the IPW estimates' high variability reflects the "genuine uncertainty" about the data (Seaman & White, 2011).

There are several ways to handle the large weights and thus stabilize the estimator. In this paper, we only discuss the weight truncation method. In weight truncation, a maximum bound is selected. Then, all the weights beyond this value will be set equal to it. As a result, we ensure a few individual observations do not excessively influence the analysis model.

5 Simulation Studies

5.1 Data Generation

We designed a series of numerical simulations to compare and contrast different estimators' performances under varied conditions. Each instance of the simulated data set is generated in two steps according to the ODS study design. In the first step, we generate the full data sets from a response model as follows:

$$Y_i = \beta_0 + \beta_1 X_i + \beta_2 Z_i + \epsilon_i$$

where $i = 1, \dots, N$, $X_i \stackrel{IID}{\sim} N(0, \phi^2)$, $Z_i \stackrel{IID}{\sim} \text{Bernoulli}(\theta)$, and $\epsilon_i \stackrel{IID}{\sim} N(0, \sigma^2)$. In the second step, based on the data collection design in Section 3, Z_i will not be observed $\forall i \in \bar{S}$. The completeness variable R_i , which indicates the observation of Z_i , is generated for each individual as follows:

$$\begin{aligned} Y_i^* &= \gamma_0 + \gamma_1 X_i + \gamma_2 Y_i + u_i \\ R_i &= \mathbb{I}(Y_i^* > 0) \end{aligned}$$

where u_i 's are independent and identically distributed random errors that follow $\text{Logistic}(0, s)$ and s is the scale parameter of the distribution. For each individual, Y_i^* is an auxiliary latent variable that controls the observation of Z_i . It can be shown that $\text{logit} \Pr(R_i = 1 | X_i, Y_i)$ is a linear combination of X_i , Y_i , and an intercept constant. The conditional distributions of the completeness variables, R_i 's, are Bernoulli distributions with varied success rates that depend on observed variables' values. Based on R_i 's value, we decide whether or not to include the individual's Z_i in the final simulated data based on the ODS design.

The parameter that defines the full data set is $\Theta = (\theta, \phi^2, \sigma^2, \beta_0, \beta_1, \beta_2)$. $\beta = (\beta_0, \beta_1, \beta_2)$ is parameter of interest, and $(\theta, \phi^2, \sigma^2)$ are the nuisance parameters that assist to define the predictor variables' distributions. s and $\gamma = (\gamma_0, \gamma_1, \gamma_2)$ are the parameters that define the missingness mechanism. We use R_{theory}^2 and R_{latent}^2 to represent the response model's and the latent model's coefficients of determination, respectively. We use R_{theory}^2 and R_{latent}^2 to determine σ^2 and s , respectively, while holding the other parameters fixed. Thus, we are able to show the estimators' performance under varied levels of the signal-to-noise ratio in the data. For each replication of simulated data set, the estimate for β is computed using the full data set, the CC analysis, the IPW approach, and the ML approach.

5.2 Simulation Result

We generate 5,000 replications. We set $N = 100$, $(\theta, \phi, \beta_0, \beta_1, \beta_2) = (0.5, 5, 0, 0.5, 2)$, $(\gamma_1, \gamma_2, \gamma_3) = (1, 1, 1)$, and $R_{\text{latent}}^2 = 0.60$. We choose two values for R_{theory}^2 : 0.30 and 0.90. We assume that the variables follow the same linear response model defined in Section 5.1. All the estimators rely on this assumption. In Tables 1 and 2, the first column lists the results of estimating β_0 , β_1 , and β_2 using the least-squares estimator with all the observations in the full data set. The other three columns list the results of estimating the parameter using different analysis methods with the ODS-designed data set. The second column lists the results of the CC analysis, the third column lists the results of the IPW method, and the fourth column lists the results of the ML method using the EM algorithm. We also examined the coverage rate of each estimator's 95 percent confidence interval (CI). It's defined as the proportion of CIs capturing the true parameter over the 5,000 iterations. For a 95 percent CI, ideally, we should expect the coverage rate falls in $[0.936, 0.964]$ (Zhao & Lipsitz, 1992).

The results in Table 1 allow a discussion about how the signal-to-noise ratio in the response model influences the various estimators' performance. The ML estimator performs well for both R_{theory}^2 values. Its results are the closest to the results from the least-squares estimator using the full data set, which is the "best" estimator. On the other hand, the IPW estimator

is the least stable among the three. When $R_{\text{theory}}^2 = 0.30$, it tends to outperform CC estimator as it is less biased and has higher coverage for β_0 and β_1 . However, When $R_{\text{theory}}^2 = 0.90$, the CC estimator outperforms the IPW estimator since both methods are unbiased and the CC estimator has higher coverage rates for all three parameters. In this scenario, it has an outstanding performance even in an ODS design, seemingly contradicting our prior knowledge in Section 4.1. Nevertheless, In fact, the high R_{theory}^2 conceals the impact of the missingness because, in this case, Y has an extremely strong relationship with X and Z . With a high R_{theory}^2 , even a few data records can generate valid estimates.

We can obtain the same observations from the estimates' density plots in Figure 1. For both R_{theory}^2 values, the ML estimates' distributions mix well with the full data least-squares estimates'. CC analysis tends to be the most biased method since the estimates' distributions deviate the most from the full data least-squares estimates'. IPW tends to be the most unstable method since the estimates' distributions are the most heavy-tailed.

Estimation of the response model parameters when the $R_{\text{theory}}^2 = 0.30$ or 0.90

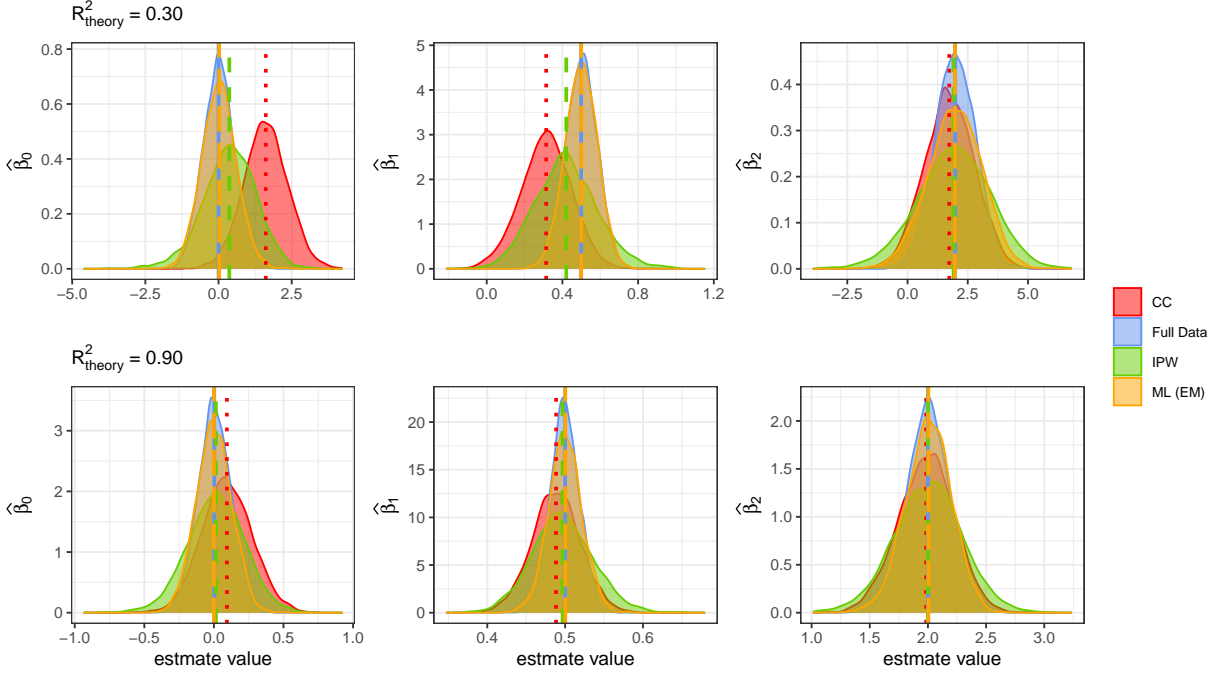


Figure 1: Density plots of the simulation results for estimating the parameters in the linear model $Y_i = \beta_0 + \beta_1 X_i + \beta_2 Z_i + \epsilon_i$, where $i = 1, \dots, 100$, $X_i \stackrel{IID}{\sim} N(0, \phi^2)$, $Z_i \stackrel{IID}{\sim} \text{Bernoulli}(\theta)$, and $\epsilon_i \stackrel{IID}{\sim} N(0, \sigma^2)$. The true parameters values are $(\theta, \phi, \beta_0, \beta_1, \beta_2) = (0.4, 5, 0, 0.5, 2)$, $(\gamma_0, \gamma_1, \gamma_2) = (1, 1, 1)$, $R_{\text{latent}}^2 = 0.60$ and $R_{\text{theory}}^2 = 0.30$ or 0.90 . The vertical lines represent the sample mean of the estimates over 5,000 replications.

Table 1: Simulation results of estimating the parameters in the linear model $Y_i = \beta_0 + \beta_1 X_i + \beta_2 Z_i + \epsilon_i$, where $i = 1, \dots, 100$, $X_i \stackrel{IID}{\sim} N(0, \phi^2)$, $Z_i \stackrel{IID}{\sim} \text{Bernoulli}(\theta)$, and $\epsilon_i \stackrel{IID}{\sim} N(0, \sigma^2)$. The true parameters values are $(\theta, \phi, \beta_0, \beta_1, \beta_2) = (0.4, 5, 0, 0.5, 2)$, $(\gamma_0, \gamma_1, \gamma_2) = (1, 1, 1)$, $R_{\text{latent}}^2 = 0.60$ and $R_{\text{theory}}^2 = 0.30$ or 0.90 . Results in the first column are calculated using the full data set. Results in the last three columns are based on the ODS data set. They are from the complete case analysis, from IPW with 20 as the truncation bound, and from the maximum likelihood method using the EM algorithm. In each row group, the first row contains the sample means of the estimates over 5,000 replications, the second row contains the sample standard deviation, and the third row contains the coverage rate, the proportion of the 95 percent CIs capturing the true parameter

R^2_{theory}	Full Data	Analysis Methods		
		CC	IPW	ML (EM)
$\beta_0 = 0$				
0.3	0.000	1.615	0.368	0.021
	0.526	0.753	0.929	0.595
	0.951	0.430	0.823	0.951
0.9	0.003	0.093	0.015	0.000
	0.115	0.181	0.206	0.128
	0.949	0.913	0.855	0.948
$\beta_1 = 0.5$				
0.3	0.499	0.314	0.420	0.499
	0.084	0.133	0.167	0.085
	0.948	0.673	0.781	0.950
0.9	0.500	0.488	0.497	0.500
	0.018	0.031	0.039	0.022
	0.948	0.925	0.810	0.948
$\beta_2 = 2$				
0.3	1.980	1.732	1.911	1.970
	0.831	1.034	1.523	1.126
	0.953	0.943	0.845	0.950
0.9	2.000	1.985	2.001	2.004
	0.186	0.240	0.296	0.199
	0.945	0.947	0.887	0.948

5.3 Model Misspecification's Influence on Estimators

We augmented the above simulations to assess response model misspecification's impact on the estimators. The same linear analysis model is used for each estimator. However, to test estimators' performance under the circumstance of mean model misspecification, we altered

the response model in data generation:

$$Y_i = \beta_0 + \beta_1 X_i^\bullet + \beta_2 Z_i + \epsilon_i$$

where X_i^\bullet is a nonlinear function of X_i . Thus, we misspecify the model if we assume Y has a linear relationship with X . The goal is to test whether the estimators yield reliable estimates for β_2 even if we make an incorrect assumption about the relationship between Y and X . We have to notice that the only thing that changes in this setting is the relation between X and Y in data generation. Both variables are collected in the final simulated data set. X^\bullet is a latent auxiliary variable. It implies we will have a model misspecification issue if we use the original response mean model. To control the degree of misspecification, we define X_i^\bullet as follows.

$$X_i^\bullet = \kappa \cdot X_i(X_i - \phi)(X_i + \phi) + X_i$$

where ϕ is the population standard deviation of X_i and κ is the misspecification level. We can use any value in the place of ϕ , but using ϕ ensures the simulated data are scattered within a reasonable range. Here, κ is used to quantify the level of misspecification. If κ is a small value, the nonlinear portion in X_i^\bullet is negligible. Then, the simulated data set in the model misspecification setting will be similar to the one generated using the original setting, and the analysis results will generally be the same. However, if κ is a large value, the simulated data and analysis result will be significantly different. Therefore, we compare the estimators' performance under varied κ .

When the analysis model is misspecified, all estimators tend to be unstable and yield estimates with high standard error. The high variability leads to a high coverage rate for $\hat{\beta}_2$ in all three methods, but the value is meaningless. The estimates can have an enormous bias, but the even larger standard error helps to construct wide CIs to cover the true value. Therefore, we utilize mean squared error (MSE) as the metric to compare estimators' efficiency. MSE captures both the bias and variance of the estimate. A lower MSE indicates better efficiency.

We generate 5,000 replications of random full data sets and their corresponding ODS-designed data sets for the model misspecification scenario. We set $N = 100$, $(\theta, \phi, \beta_0, \beta_1, \beta_2) = (0.5, 5, 0, 0.5, 2)$, $(\gamma_1, \gamma_2, \gamma_3) = (1, 1, 1)$, and $R_{\text{latent}}^2 = 0.60$. We choose two values for R_{theory}^2 : 0.30 and 0.90, and three values for κ : 0.01, 0.1, and 1. For each pair of R_{theory}^2 and κ , we perform a simulation study with the above setting. The goal of interest is to study if the incorrect assumption of the relationship between X and Y can affect the estimation of β_2 . Thus, we compare the estimators only with respect to $\hat{\beta}_2$. In Tables 3 and 4, the first column lists the results of estimating β_2 using the least-squares estimator with all the observations in the full data set. The other three columns list the results of estimating the parameter using different analysis methods with the ODS-designed data set. The second column lists the results of the CC analysis, the third column lists the results of the IPW method, and the fourth column lists the results of the ML method using the EM algorithm. We calculate the sample mean of $\text{MSE}(\hat{\beta}_2)$ over the 5,000 iteration. It can easily proven that $\text{MSE}(\hat{\beta}_2) = \text{Bias}^2(\hat{\beta}_2) + \text{Var}(\hat{\beta}_2)$.

Table 2: Simulation results of estimating only β_2 when $R^2_{\text{theory}} = 0.30$ or 0.90 , where both model well-specified scenario and model misspecified scenarios are included. The value of the parameter is 2. The four columns list the results from different analysis methods. Results in the first column are calculated using the full data set. Results in the last three columns are based on the ODS data set. They are from the complete case analysis, from IPW with 20 as the truncation bound, and from the maximum likelihood method using the EM algorithm. In each row group, the first row contains the sample means of $\hat{\beta}_2$ over 5,000 replications, the second row contains the of $\text{MSE}(\hat{\beta}_2)$.

$\beta_2 = 2$				
R^2_{theory}	Full Data	Analysis Methods		
		CC	IPW	ML (EM)
Correctly Specified				
0.3	1.980	1.732	1.911	1.970
	0.691	1.141	2.326	1.269
0.9	2.000	1.984	2.001	2.003
	0.034	0.058	0.087	0.040
Misspecified, $\kappa = 0.01$				
0.3	1.984	1.802	1.972	1.964
	0.805	1.360	2.306	1.440
0.9	1.996	2.021	1.999	1.902
	0.131	0.208	0.179	0.178
Misspecified, $\kappa = 0.1$				
0.3	1.968	1.885	1.991	1.713
	10.006	17.213	10.407	14.950
0.9	2.014	2.036	2.073	1.888
	9.542	16.733	8.796	14.386
Misspecified, $\kappa = 1$				
0.3	2.413	2.816	2.853	0.760
	959.492	1832.761	932.155	1521.635
0.9	2.261	2.706	2.585	0.513
	972.835	1946.229	985.904	1594.518

The results in Tables 2 allow a discussion about how the misspecification influences various estimators from two aspects: bias, efficiency. When $R^2_{\text{theory}} = 0.30$ and $\kappa = 0.01$, the misspecification has negligible effect on three estimators' performance. The CC estimator gives the most biased estimate, but it tends to be more efficient than the other estimators under the MSE criterion. IPW estimator is less biased, but it is more unstable and inefficient than other estimators. ML estimator is less biased than the CC estimator, and it is more stable and efficient than the IPW estimator. However, when the degree of misspecification becomes more significant, $\kappa = 0.1$, all three estimators receive different levels of influence. The IPW estimator becomes the most unbiased, stable, and efficient estimator among the

three. On the contrary, the CC estimator becomes more unstable than IPW, and the ML estimator becomes the most biased and inefficient. When $\kappa = 1$, all three estimators become highly biased and inefficient, but the IPW estimator remains the most unbiased and efficient among them. When $R^2_{\text{theory}} = 0.90$ and $\kappa = 0.01$, all three estimators are unbiased and efficient. Similarly, all estimators receive varied levels of impact as the response model becomes more misspecified. The IPW estimator receives the least effect, while the ML estimator is the most vulnerable to model misspecification.

6 Discussion

Our simulation results corroborate the results from the previous work. In an ODS study, CC analysis can be reliable only when the response variable has a strong relationship with the predictor variables. The ML method achieves its ideal performance when our assumption about the response model is valid. Once the assumption is incorrect, the estimator can be highly biased and inefficient since the EM algorithm is sensitive. Moreover, no software package supports the EM algorithm by default, and the development of the algorithm requires a considerable amount of time and effort. On the other hand, the IPW approach, although not ideally unbiased and efficient as the ML method, is easy to implement in many software packages. Therefore, from our limited experience, we make the following recommendations. Unless there is evidence that the response variable has a perfect relationship with the predictor variables, CC analysis does not correctly address the missingness. Instead, IPW can be used as an improvised method. After thoroughly learning the data set, we can use the ML method and develop an EM algorithm to calculate the estimates.

Several extensions of the design and methods are worth mentioning. The first extension is to compare and contrast IPW and EM estimators' efficiency using bootstrapping. The case resampling process must be stratified to maintain the degree of missingness in each bootstrapped data set. Another extension is to explore the poor coverage rate of the IPW estimator. In our simulation results, the 95 percent CI coverage rates of the IPW estimates do not fall into $[0.936, 0.964]$. However, Weaver et al. (2005) have the IPW estimators achieve reasonable coverage rates in their ODS design. The reason for the difference is that the IPW estimates standard errors we use are directly from the software package by default, which is commonly used in practice. However, these values are not theoretically correct. Therefore, we need to develop our own functions to compute the theoretical standard errors.

Appendix

In our ODS design, we have the following setup

$$\begin{aligned} y_i | x_i, z_i &\overset{Ind}{\sim} N(\beta_0 + \beta_1 x_i + \beta_2 z_i, \sigma^2) \\ x_i &\overset{IID}{\sim} N(0, \phi^2) \\ z_i &\overset{IID}{\sim} \text{Bernoulli}(\theta) \\ R_i &= \mathbb{I}(z_i \text{ is observed}) \end{aligned}$$

The observed likelihood is

$$L_{Obs}(\Theta) = \prod_{i=1}^N \left[f(y_i, x_i; \Theta) f(z_i | y_i, x_i; \Theta) \right]^{R_i} \left[f(y_i, x_i; \Theta) \right]^{1-R_i}$$

We also know the complete likelihood is

$$\begin{aligned} L(\Theta) &= f(\mathbf{y}, \mathbf{x}, \mathbf{z}_S, \mathbf{z}_{\bar{S}}; \Theta) \\ &= \prod_{i=1}^N f(z_i | y_i, x_i; \Theta) f(y_i, x_i; \Theta) \end{aligned}$$

Then, the conditional likelihood of $\mathbf{z}_{\bar{S}}$ given all the observed data is

$$\begin{aligned} f(\mathbf{z}_{\bar{S}} | \mathbf{y}, \mathbf{x}, \mathbf{z}_S; \Theta) &= \frac{f(\mathbf{y}, \mathbf{x}, \mathbf{z}_S, \mathbf{z}_{\bar{S}}; \Theta)}{f(\mathbf{y}, \mathbf{x}, \mathbf{z}_S; \Theta)} \\ &= \frac{L(\Theta)}{L_{Obs}(\Theta)} \\ &= \prod_{i=1}^N \left[f(z_i | y_i, x_i) \right]^{1-R_i} \end{aligned}$$

Note that $L(\Theta)$ can also be expressed as $f(\mathbf{y} | \mathbf{x}, \mathbf{z}) f(\mathbf{x}) f(\mathbf{z})$.

Expectation Step

By definition,

$$Q(\Theta | \hat{\Theta}^{(m)}) = \mathbb{E}[\log L(\Theta)]$$

where the expectation is over $f(\mathbf{z}_{\bar{S}} | \mathbf{y}, \mathbf{x}, \mathbf{z}_S; \Theta)$. Thus, due to the independence among all observations, the Q function can be expressed as follows:

$$\begin{aligned} Q(\Theta | \hat{\Theta}^{(m)}) &= \sum_{i=1}^N \left(R_i \log f(y_i | x_i, z_i) + (1 - R_i) \cdot \mathbb{E}[\log f(y_i | x_i, z_i)] \right. \\ &\quad \left. + \log f(x_i) + R_i \log f(z_i) + (1 - R_i) \cdot \mathbb{E}[\log f(z_i)] \right) \end{aligned}$$

where x_i 's distribution parameters does not affect the estimation for β so that $\log f(x_i)$ can be ignored in the future calculation.

We can prove that

$$f(y_i, x_i) = f(y_i | x_i, z_i = 0) f(x_i) (1 - \theta) + f(y_i | x_i, z_i = 1) f(x_i) \theta$$

and

$$z_i | y_i, x_i \stackrel{Ind}{\sim} \text{Bernoulli}(\phi'_i), \forall i \in \bar{S}$$

$$\text{where } \phi'_i = \frac{f(y_i | x_i, z_i = 1) \theta}{f(y_i | x_i, z_i = 1) \theta + f(y_i | x_i, z_i = 0) (1 - \theta)}$$

Therefore, we can express the Q function as follows:

$$Q(\Theta|\hat{\Theta}^{(m)}) = \sum_{i=1}^N \left[R_i \log f(y_i|x_i, z_i) + R_i \log f(z_i) \right. \\ \left. + (1 - R_i) \log f(y_i|x_i, z_i = 1)\phi'_i + (1 - R_i) \log f(y_i|x_i, z_i = 0)(1 - \phi'_i) \right. \\ \left. + (1 - R_i) \left(\phi'_i \log(\theta) + (1 - \phi'_i) \log(1 - \theta) \right) \right]$$

which is iteratively maximized during the maximization step.

Maximization Step

In this step, we iteratively maximize the Q function by letting

$$\hat{\Theta}^{(m+1)} = \arg \max Q(\Theta|\hat{\Theta}^{(m)})$$

$Q(\Theta|\hat{\Theta}^{(m)})$ can be decomposed as follows:

$$Q(\Theta|\hat{\Theta}^{(m)}) = Q_\theta + Q_{\beta, \sigma^2}$$

$$Q_\theta = \sum_{i=1}^N R_i \log f(z_i) + \sum_{i=1}^N (1 - R_i) \left(\phi'_i \log(\theta) + (1 - \phi'_i) \log(1 - \theta) \right)$$

$$Q_{\beta, \sigma^2} = \sum_{i=1}^N R_i \log f(y_i|x_i, z_i) + \sum_{i=1}^N \log f(y_i|x_i, z_i = 0)(1 - \phi'_i) \\ + \sum_{i=1}^N (1 - R_i) \log f(y_i|x_i, z_i = 1)\phi'_i$$

Then, we can maximize the estimates element-wisely.

For θ , we can directly calculate the updated estimate by setting $\frac{d}{d\theta} Q_\theta = 0$. For β and σ^2 , we can use the default weighted least-squares estimation in many software packages. It can be derived that

$$Q_{\beta, \sigma^2} = -\frac{n}{2} \log(2\pi\sigma^2) - \frac{S(\beta)}{2\sigma^2}$$

where

$$S(\beta) = \sum_{i=1}^N \left(R_i (y_i - \beta_0 - \beta_1 x_i - \beta_2 z_i)^2 + (1 - R_i) \left((y_i - \beta_0 - \beta_1 x_i - \beta_2)^2 \phi'_i \right. \right. \\ \left. \left. + (y_i - \beta_0 - \beta_1 x_i)^2 (1 - \phi'_i) \right) \right)$$

The updated estimate for β is calculated by

$$\arg \max_{\beta} Q_{\beta, \sigma^2} = \arg \min_{\beta} S(\beta)$$

which is essentially the weighted least-square estimator of an auxiliary data set. The auxiliary data set consists of three components: the complete cases, the incomplete cases with z_i 's imputed as 1, and the incomplete cases with z_i 's imputed as 0. After obtaining the updated estimate for β , we can calculate the updated estimate of σ^2 by setting $\frac{d}{d\theta} Q_{\beta, \sigma^2} = 0$. It is noteworthy that we also need to update ϕ_i 's during each iteration.

The algorithm terminates when the minimum change in the estimates, $\min(d\hat{\Theta})$, is less than a tolerance value.

Reference

- Schafer, Joseph L, and John W Graham. 2002. "Missing Data: Our View of the State of the Art." *Psychological Methods* 7 (2): 147.
- Seaman, Shaun R, and Ian R White. 2013. "Review of Inverse Probability Weighting for Dealing with Missing Data." *Statistical Methods in Medical Research* 22 (3): 278–95.
- Weaver, Mark A, and Haibo Zhou. 2005. "An Estimated Likelihood Method for Continuous Outcome Regression Models with Outcome-Dependent Sampling." *Journal of the American Statistical Association* 100 (470): 459–69.
- White, Ian R, and John B Carlin. 2010. "Bias and Efficiency of Multiple Imputation Compared with Complete-Case Analysis for Missing Covariate Values." *Statistics in Medicine* 29 (28): 2920–31.
- Zhao, LP, and S Lipsitz. 1992. "Designs and Analysis of Two-Stage Studies." *Statistics in Medicine* 11 (6): 769–82.